

CSCB09 Software Tools and Systems Programming

IO Multiplexing // Threads

Marcelo Ponce

Winter 2025

Department of Computer and Mathematical Sciences - UTSC

Today's class

Today we will discuss the following topics:

- IO Multiplexing
- Threads

IO Multiplexing

- One basic concept of Linux systems is the rule that everything in Unix/Linux is a file.
- Each process has a table of file descriptors that point to files, sockets, devices and other operating system objects.
- Not so surprising then is the fact, that processes will have to deal with a large number of files, i.e. FDs.
- A simple solution would be to create a process (or thread) to deal with them...

- One basic concept of Linux systems is the rule that everything in Unix/Linux is a file.
- Each process has a table of file descriptors that point to files, sockets, devices and other operating system objects.
- Not so surprising then is the fact, that processes will have to deal with a large number of files, i.e. FDs.
- A simple solution would be to create a process (or thread) to deal with them... For small numbers of FDs this may work, but for large ones is a bad idea!

- One basic concept of Linux systems is the rule that everything in Unix/Linux is a file.
- Each process has a table of file descriptors that point to files, sockets, devices and other operating system objects.
- Not so surprising then is the fact, that processes will have to deal with a large number of files, i.e. FDs.
- A simple solution would be to create a process (or thread) to deal with them... For small numbers of FDs this may work, but for large ones is a bad idea!

The solution to this, is to use a kernel mechanism for *polling* –check the status– over a set of FDs:

Multiplexing

a system or signal involving simultaneous transmission of several messages along a single channel of communication.

- One basic concept of Linux systems is the rule that everything in Unix/Linux is a file.
- Each process has a table of file descriptors that point to files, sockets, devices and other operating system objects.
- Not so surprising then is the fact, that processes will have to deal with a large number of files, i.e. FDs.
- A simple solution would be to create a process (or thread) to deal with them... For small numbers of FDs this may work, but for large ones is a bad idea!

The solution to this, is to use a kernel mechanism for *polling* –check the status– over a set of FDs:

Multiplexing


a system or signal involving simultaneous transmission of several messages along a single channel of communication.

There are three main options in Linux: `select`, `poll`, `epoll`

- Typical system that works with many IO sources has an initialization phase and then enter some kind of standby mode – wait for any client to send request and response it
- A tentative “solution” is to create a thread (or process) for each client, block on read until a request is sent and write a response.
- This may work with a small amount of clients, but if we want to scale it to hundred of clients, creating a thread for each client is a bad idea


- Typical system that works with many IO sources has an initialization phase and then enter some kind of standby mode – wait for any client to send request and response it
- A tentative “solution” is to create a thread (or process) for each client, block on read until a request is sent and write a response.
- This may work with a small amount of clients, but if we want to scale it to hundred of clients, creating a thread for each client is a bad idea

The solution is to use a kernel mechanism for polling over a set of file descriptors.



I/O Models

- Blocking I/O
- Non-blocking I/O
- I/O Multiplexing (select/poll/...)
- Signal driven I/O (SIGIO)
- Asynchronous I/O



select system call

The `select()` system call provides a mechanism for implementing synchronous multiplexing I/O

```
int select(int nfd, fd_set *readfds, fd_set *writefds, fd_set *exceptfds,
           struct timeval *timeout);
```

A call to `select()` will **block** until the given FDs are ready to perform I/O, or until an optionally specified timeout has elapsed.

poll system call

Unlike `select()`, with its inefficient three bitmask-based sets of file descriptors, `poll()` employs a single array of `nfds` `pollfd` structures

```
int poll (struct pollfd *fds, unsigned int nfds, int timeout);
```

```
struct pollfd {  
    int fd;  
    short events;  
    short revents;  
};
```

epoll* system calls help us to create and manage the context in the kernel

- create a context in the kernel using `epoll_create`
- add and remove file descriptors to/from the context using `epoll_ctl`
- wait for events in the context using `epoll_wait`

select **VS** poll **VS** epoll

select

- FD sets statically sized
- FD sets reconstructed on return
- timeout parameter undefined on return
- more portable
- pselect

poll

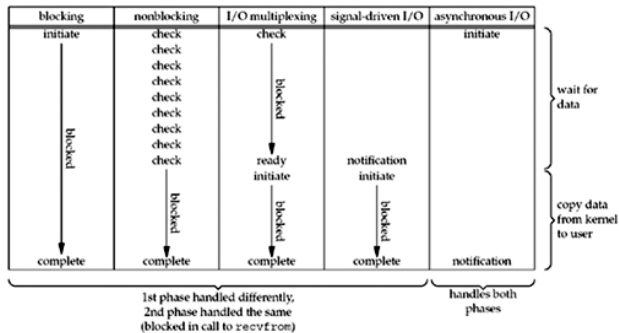

- does not require user to calculate max nbr of FDs
- more efficient for large-valued FDs
- separates input (events) from the output (revents)
- more portable

epoll

- can add/remove FDs while waiting
- better performance: $\mathcal{O}(1)$ instead of $\mathcal{O}(n)$
- can behave as level triggered or edge triggered
- linux specific, so non portable

Libevent Benchmark

One Active Connection



Concurrency

- The two key concepts driving computer systems and applications are
 - **communication**: the conveying of information from one entity to another
 - **concurrency**: the sharing of resources in the same time frame
- Concurrency can exist in a single processor as well as in a multiprocessor system
- Managing concurrency is difficult, as execution behaviour is not always reproducible

Concurrency – example

```
#!/usr/bin/sh
count=1
while [ $count -le 20 ]
do
    echo -n "a"
    count=`expr $count + 1`
done
```

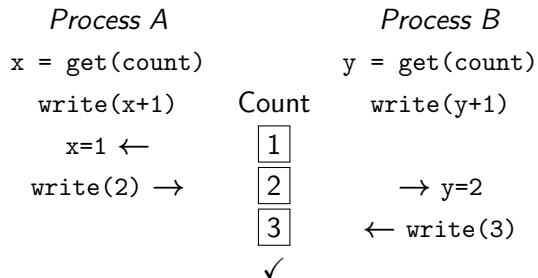
```
#!/usr/bin/sh
count=1
while [ $count -le 20 ]
do
    echo -n "b"
    count=`expr $count + 1`
done
```

- When run sequentially (a; b) output is sequential.
- When run concurrently (a&; b&) output is interspersed and different from run to run.

Race Conditions

- A **Race Condition** occurs when multiple processes are trying to do something with shared data and the final outcome depends on the order in which the processes run.
- E.g., If any code after a fork depends on whether the parent or child runs first.
- A parent process can call `wait()` to wait for termination (may block)
- A child process can wait for parent to terminate by polling (wasteful)
(How would you do this?)
- One standard solution is to use signals.

Race Condition – Example



Race Condition – Example

Process A
x = get(count)
write(x+1)
x=1 ←
write(2) →

Count

1

2

3

✓

Process B
y = get(count)
write(y+1)
→ y=2
← write(3)

Process A
x = get(count)
write(x+1)
x=1 ←

write(2) →

Count

1

2

3

2


Process B
y = get(count)
write(y+1)
→ y=1
← write(2)
→ y=2
← write(3)

Threads

- Processes are expensive to create.
- It takes quite a bit of time to switch between processes
- Communication between processes must be done through an external structure
 - files, pipes, shared memory
- Synchronizing between processes is cumbersome.

- Processes are expensive to create.
- It takes quite a bit of time to switch between processes
- Communication between processes must be done through an external structure
 - files, pipes, shared memory
- Synchronizing between processes is cumbersome.
- Is there another model that will solve these problems?

Processes



- Each process has its own
 - program counter
 - stack
 - stack pointer
 - address space
- Processes may share
 - open files
 - pipes

What is a process?

- OS abstraction for execution
- Running instance of a program

What is a process?

- OS abstraction for execution
- Running instance of a program
- Components of a process:
 - Address space
 - Code and data
 - Stack
 - Program Counter (PC)
 - Set of registers
 - Set of OS resources: open files, network connections...

Rethinking Processes

- What is similar in cooperating processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)


Rethinking Processes

- What is similar in cooperating processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)
- What don't they share?
 - Each has its own execution state: PC, SP, and registers

Rethinking Processes

- What is similar in cooperating processes?
 - They all share the same code and data (address space)
 - They all share the same privileges
 - They all share the same resources (files, sockets, etc.)
- What don't they share?
 - Each has its own execution state: PC, SP, and registers
- **Key idea:** Why don't we separate the concept of a process from its execution state?
 - **Process:** address space, privileges, resources, etc.
 - **Execution state:** PC, SP, registers
- *Execution state* also called **thread of control**, or **thread**

Threads



- Each thread has its own
 - program counter
 - stack
 - stack pointer
- Threads share
 - address space
 - variables
 - code
 - open files

What is a Thread?

- A **thread** is a single **control flow** through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with *multiple control flows* is **multithreaded**

What is a Thread?

Control Flow

- A **thread** is a single **control flow** through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with *multiple control flows* is **multithreaded**

What is a Thread?

- A **thread** is a single **control flow** through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with *multiple control flows* is **multithreaded**

Control Flow

- Control includes all of the values that select which instructions in a program are executed.

What is a Thread?

- A **thread** is a single **control flow** through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with *multiple control flows* is **multithreaded**

Control Flow

- Control includes all of the values that select which instructions in a program are executed.
- Control flow, then, is the sequence of instructions being executed.

What is a Thread?

- A **thread** is a single **control flow** through a program
 - What is a “control flow”?
 - How is control flow represented?
- A program with *multiple control flows* is **multithreaded**

Control Flow

- Control includes all of the values that select which instructions in a program are executed.
- Control flow, then, is the sequence of instructions being executed.
- The hardware uses the program counter (PC) and stack to make control flow decisions.

Threads – Advantages

- Communication between threads is cheap
 - they can share variables!
- Threads are “lightweight”
 - faster to create
 - faster to switch between

Producer/Consumer Problem

- Simple example:

```
who | wc -l
```

- Both the writing process (`who`) and the reading process (`wc`) of a pipeline execute concurrently.
- A pipe is usually implemented as an internal OS buffer.
- It is a resource that is **concurrently** accessed by the reader and the writer, so it must be managed carefully.

Producer/Consumer

- **consumer** should be blocked when buffer is empty
- **producer** should be blocked when buffer is full
- producer and consumer should run independently as far as buffer capacity and contents permit
- producer and consumer should never be updating the buffer at the same instant (otherwise data integrity cannot be guaranteed)
- producer/consumer is a harder problem if there are more than one consumer and/or more than one producer.

Pthreads

- POSIX threads (pthreads) is the most commonly used thread package on Unix/Linux

- POSIX threads (pthreads) is the most commonly used thread package on Unix/Linux

POSIX: Portable Operating System-(ix)

a set of formal descriptions that provide a standard for the design of operating systems, especially ones which are compatible with Unix.

```
int pthread_create(pthread_t *tid, pthread_attr_t *attr,  
void *(*func)(void*), void *arg);
```

- tid uniquely identifies a thread within a process and is returned by the function
- attr sets attributes such as priority, initial stack size
 - can be specified as NULL to get defaults
- func – the function to call to start the thread
 - accepts one void *argument, returns void *
- arg is the argument to func
- returns 0 if successful, a positive error code if not
- does not set errno but returns compatible error codes
- can use strerror() to print error messages

```
int pthread_join(pthread_t tid, void **status)
```

- tid – the tid of the thread to wait for
 - cannot wait for any thread (as in wait())
- status, if not NULL returns the void * returned by the thread when it terminates.
- a thread can terminate by
 - returning from func
 - the main() function exiting
 - pthread_exit()

More `pthread` Functions

More `pthread` Functions

- `void pthread_exit(void *status)`
 - a second way to exit, returns status explicitly
 - status must not point to an object local to the thread, as these disappear when the thread terminates

More pthread Functions

- `void pthread_exit(void *status)`
 - a second way to exit, returns status explicitly
 - status must not point to an object local to the thread, as these disappear when the thread terminates
- `void pthread_detach(pthread_tid)`
 - if a thread is detached its termination cannot be tracked with `pthread_join()`
 - it becomes a daemon thread

More pthread Functions

- `void pthread_exit(void *status)`
 - a second way to exit, returns status explicitly
 - status must not point to an object local to the thread, as these disappear when the thread terminates
- `void pthread_detach(pthread_tid)`
 - if a thread is detached its termination cannot be tracked with `pthread_join()`
 - it becomes a daemon thread
- `pthread_t pthread_self(void)`
 - returns the thread ID of the thread which called it
 - often see `pthread_detach(pthread_self())`

Passing Arguments to Threads

```
pthread_t thread_ID; int fd, result;
fd = open("afile", "r");
result = pthread_create(&thread_ID, NULL, myThreadFcn, (void *)&fd);

if (result != 0)
    printf("Error: %s\n", strerror(result));
```

- We can pass any variable (including a structure or array) to our thread function.
- It assumes the thread function knows what type it is.
- This example is bad if the main thread alters fd later.

Passing Arguments to Threads

```
pthread_t thread_ID; int fd, result;
fd = open("afile", "r");
result = pthread_create(&thread_ID, NULL, myThreadFcn, (void *)&fd);

if (result != 0)
    printf("Error: %s\n", strerror(result));
```

- We can pass any variable (including a structure or array) to our thread function.
- It assumes the thread function knows what type it is.
- This example is bad if the main thread alters fd later.

Solution,

- Use malloc() to create memory for the variable
 - initialize variable's value
 - pass pointer to new memory via pthread_create()
 - thread function releases memory when done.

Example

```
typedef struct myArg { int fd;
    char name[25];
} MyArg;

int result;
pthread_t thread_ID;

MyArg *p = (MyArg *)malloc(sizeof(MyArg));
p->fd = fd; /* assumes fd is defined */
strncpy(p->name, "CSCB09", 7);
result = pthread_create(&threadID, NULL, myThreadFcn, (void *)p);

void *myThreadFcn(void *p) {
    MyArg *theArg = (MyArg *) p;
    write(theArg->fd, theArg->name, 7);
    close(theArg->fd);
    free(theArg);
    return NULL;
}
```

* Needs to include `-pthread` flag in compilation!

Thread-Safe Functions

- Not all functions can be called from threads
 - many use global/static variables
 - new versions of UNIX have thread-safe replacements like `strtok_r()`

Thread-Safe Functions

- Not all functions can be called from threads
 - many use global/static variables
 - new versions of UNIX have thread-safe replacements like `strtok_r()`
- Safe:
`ctime_r()`, `gmtime_r()`, `localtime_r()`, `rand_r()`, `strtok_r()`



Thread-Safe Functions

- Not all functions can be called from threads
 - many use global/static variables
 - new versions of UNIX have thread-safe replacements like `strtok_r()`
- Safe:
`ctime_r()`, `gmtime_r()`, `localtime_r()`, `rand_r()`, `strtok_r()`
- Not Safe:
`ctime()`, `gmtime()`, `localtime()`, `rand()`, `strtok()`, `gethostxxx()`



Thread-Safe Functions

- Not all functions can be called from threads
 - many use global/static variables
 - new versions of UNIX have thread-safe replacements like `strtok_r()`
- Safe:
`ctime_r()`, `gmtime_r()`, `localtime_r()`, `rand_r()`, `strtok_r()`
- Not Safe:
`ctime()`, `gmtime()`, `localtime()`, `rand()`, `strtok()`, `gethostxxx()`
- Could use semaphores to protect access but will generally result in poor performance.

Pthread Mutexes



Pthread Mutexes



```
int pthread_mutex_init(pthread_mutex_t *mp,  const pthread_mutexattr_t *attr);  
  
int pthread_mutex_lock(pthread_mutex_t *mp);  
int pthread_mutex_trylock(pthread_mutex_t *mp);  
int pthread_mutex_unlock(pthread_mutex_t *mp);  
int pthread_mutex_destroy(pthread_mutex_t *mp);
```

- easier to use than `semget()` and `semop()`
- only the thread that locks a mutex can unlock it
- mutexes often declared as globals

Example

```
////////////////////////////////////////  
  
pthread_mutex_t myMutex;  
int status;  
  
status = pthread_mutex_init(&myMutex, NULL);  
if(status != 0)  
    printf("Error: %s\n", strerror(status));  
  
pthread_mutex_lock(&myMutex);  
/* critical section here */  
pthread_mutex_unlock(&myMutex);  
  
status = pthread_mutex_destroy(&myMutex);  
if(status != 0)  
    printf("Error: %s\n", strerror(status));  
  
////////////////////////////////////////
```


Concurrency

- Critical Region
- Atomic instructions
- Barrier
- Higher level abstraction
OpenMP

OpenMP


OpenMP

- OMP uses explicit (not automatic) parallelism
- compiler-directive based
- fork-join model



OpenMP

- OMP uses explicit (not automatic) parallelism
- compiler-directive based
- fork-join model



```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <unistd.h>
4 #include <omp.h>
5
6 int main() {
7
8     //fork();
9     #pragma omp parallel for
10    for (int i=1; i<10; i++) {
11        printf("hello %d from thread %d\n", i, omp_get_thread_num());
12    }
13
14    return 0;
15 }
```

OpenMP Example

```
1 #include <omp.h>
2
3 main () {
4     int nthreads, tid;
5
6     /* Fork a team of threads with each thread having a private tid variable */
7     #pragma omp parallel private(tid) {
8
9         /* Obtain and print thread id */
10        tid = omp_get_thread_num();
11        printf("Hello World from thread = %d\n", tid);
12
13        /* Only master thread does this */
14        if (tid == 0) {
15            nthreads = omp_get_num_threads();
16            printf("Number of threads = %d\n", nthreads);
17        }
18
19    } /* All threads join master thread and terminate */
20 }
```

References and Further Resources

- Networking and IO Models
<https://seenaburns.com/network-io/>
- Comparison of Threading Programming Models
<https://par.nsf.gov/servlets/purl/10050480>
- POSIX-Threads and OpenMP (Shared-Memory Paradigm)
https://www.cs.uic.edu/~ajayk/c566/Presentation_POSIX_OpenMP.pdf
- POSIX Threads Programming (LLNL)
<https://hpc-tutorials.llnl.gov/posix/>

- Networking and IO Models
<https://seenaburns.com/network-io/>
- Comparison of Threading Programming Models
<https://par.nsf.gov/servlets/purl/10050480>
- POSIX-Threads and OpenMP (Shared-Memory Paradigm)
https://www.cs.uic.edu/~ajayk/c566/Presentation_POSIX_OpenMP.pdf
- POSIX Threads Programming (LLNL)
<https://hpc-tutorials.llnl.gov/posix/>

If you are interested to learn more about shared-memory (threads) and distributed-memory (processes) parallelism, take a look at this **Fall'25 CSCD71** *"Special Topics course in HPC"*