

CSCB09 Software Tools and Systems Programming

I/O & Memory Model

Marcelo Ponce

Winter 2025

Department of Computer and Mathematical Sciences - UTSC

Today's lecture

File System

File IO

Arrays & Pointers

Memory Model

File System

File System, Inodes and Blocks

- Components of the OS in charge of managing *files*.
- Interact with lower-level IO subsystems
- All Unix filesystems use two basic components to organize and store data: **blocks** and **inodes**.

File System, Inodes and Blocks

- Components of the OS in charge of managing *files*.
- Interact with lower-level IO subsystems
- All Unix filesystems use two basic components to organize and store data: **blocks** and **inodes**.

Similarly as a physical disk is organized into *sectors*, data on a filesystem is abstracted into *blocks*.

File System, Inodes and Blocks

- Components of the OS in charge of managing *files*.
- Interact with lower-level IO subsystems
- All Unix filesystems use two basic components to organize and store data: **blocks** and **inodes**.

Blocks

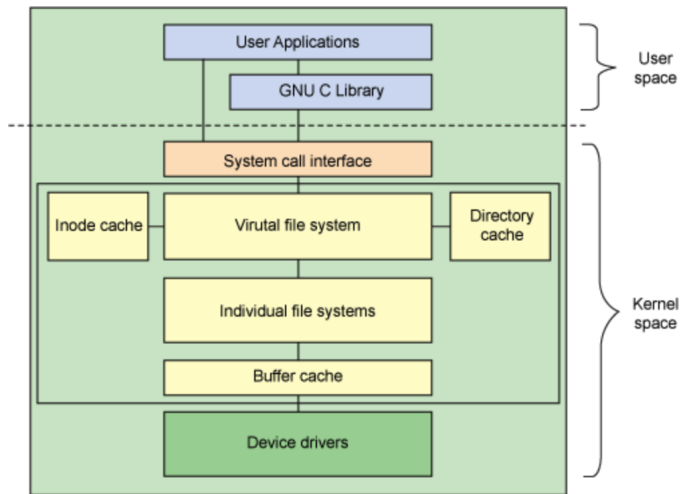
- abstractions of data on the filesystem
- have a fixed size, determined at the time the filesystem is created
- the *block size* determines how many bytes are allocated to each block
- e.g. 32-bit: 1 KB, 2 KB, or 4 KB;
64-bit: 8 KB

Similarly as a physical disk is organized into *sectors*, data on a filesystem is abstracted into *blocks*.

Inodes

- data structure used to map blocks to physical disk locations
- every file, is assigned an inode
- analogous to pointers
- when a filesystem runs out of inodes, no new files can be created until existing files are deleted

Linux file system Architecture



SRC: <https://developer.ibm.com/tutorials/l-linux-filesystem/>

File IO

Basic Input/Output Operations – IOPS

What are I/O operations, or IOPS?

- Finding a file (ls)
Check if that file exists, read metadata (file size, date stamp etc.)
- Opening a file:
Check if that file exists, see if opening the file is allowed, possibly create it, find the block that has the (first part of) the file system.
- Reading a file:
Position to the right spot, read a block, take out right part
- Writing to a file:
Check where there is space, position to that spot, write the block. Repeated if the data read/written spans multiple blocks.
- Move the file pointer (“seek”):
File system must check where on disk the data is.
- Close the file.

Basic File I/O functions (stdio.h)

```
// Open
FILE *fopen(const char *path, const char *mode);

// Input:
int fscanf(FILE *stream, const char *format, ...);
char *fgets(char *s, int size, FILE *stream);
char fgetc(FILE *stream);

// Output:
int fprintf(FILE *stream, const char *format, ...);
int fputs(const char *str, FILE *stream);

// Position cursor:
int fseek(FILE *stream, long int offset, int whence);
void rewind(FILE *stream);

// Close
int fclose(FILE *stream);
```

Some reflections about basics of file I/O and command line arguments...

Questions...?

- Why do we need to open files before reading/writing on them?
- Why do we need to close files when done with them?
- Why did we not have to open STDOUT before we could use printf?

File interfaces in C/Unix

- Two main mechanisms for managing open files:

File interfaces in C/Unix

- Two main mechanisms for managing open files:
- File *descriptors* (low-level, managed by OS):
 - Each open file is identified by a small integer
 - STDIN is 0, STDOUT is 1
 - Use for pipes, sockets (will see later what those are ...)
 - (Remember how I said in first lecture that everything is a file ..?)

File interfaces in C/Unix

- Two main mechanisms for managing open files:
- File *descriptors* (low-level, managed by OS):
 - Each open file is identified by a small integer
 - STDIN is 0, STDOUT is 1
 - Use for pipes, sockets (will see later what those are ...)
 - (Remember how I said in first lecture that everything is a file ..?)
- File *pointers* (aka streams, file handles) for regular files:
 - A C language construct for easier working with file descriptors
 - You use a pointer to a file structure (FILE *) as handle to a file.
 - The file struct contains a file descriptor and a buffer.
 - Use for regular files

Standard File Descriptors

- The operating system (OS) maintains **for each process** a table with open file descriptors: **file descriptor table**

Standard File Descriptors

- The operating system (OS) maintains **for each process** a table with open file descriptors: **file descriptor table**
- The OS opens for each process automatically three standard file descriptors:

	Name	File Descriptor	Default
Standard Input	stdin	0	keyboard
Standard Output	stdout	1	screen
Standard Error	stderr	2	screen

Standard File Descriptors

- The operating system (OS) maintains **for each process** a table with open file descriptors: **file descriptor table**
- The OS opens for each process automatically three standard file descriptors:

	Name	File Descriptor	Default
Standard Input	stdin	0	keyboard
Standard Output	stdout	1	screen
Standard Error	stderr	2	screen

- All C programs automatically have three files stream open:

```
FILE *stdin;  
FILE *stdout;  
FILE *stderr;
```

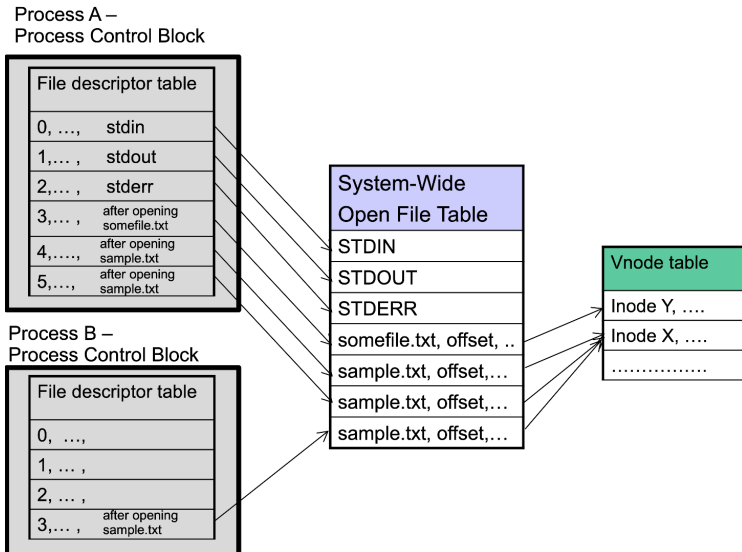
- Those are ready to use:

```
// using fprintf and a file handler  
fprintf(stdout, "Hello!\n");  
  
// identical to use printf  
printf("Hello!\n");
```

OS managing open files (in-memory data structures)

A new entry is added to the file descriptor table each time a file is opened

- Find *inode*
- Check permissions
- Initialize cursor to beginning of file



Arrays & Pointers

Arrays & Pointers

```
int x[5];  
int *y = NULL;  
  
for (int i = 0; i < 5; i++){  
    x[i] = 1;  
}
```

Arrays & Pointers

```
int x[5];
int *y = NULL;

for (int i = 0; i < 5; i++){
    x[i] = 1;
}
```

y	NULL	0x88681140
x[0]	1	0x88681144
x[1]	1	0x88681148
x[2]	1	0x8868114c
x[3]	1	0x88681150
x[4]	1	0x88681154
	?	0x88681158

If an array of ints contains 10 ints, then the array size is 40 bytes. There is nothing extra.

- In particular, size is not stored.
- No memory is reserved to hold the address where array starts.

For a pointer variable, space to hold size of an address (4 or 8 bytes) is reserved.

Arrays & Pointers

```
int x[5];  
int *y = NULL;  
y = x;  
  
for (int i = 0; i < 5; i++)  
{  
    y[i] = 1;  
}
```

Arrays & Pointers

```
int x[5];  
int *y = NULL;  
y = x;  
  
for (int i = 0; i < 5; i++)  
{  
    y[i] = 1;  
}
```

y		0x88681140
x[0]		0x88681144
x[1]		0x88681148
x[2]		0x8868114c
x[3]		0x88681150
x[4]		0x88681154
		0x88681158

Arrays & Pointers

```
int x[5];
int *y = NULL;
y = x;

for (int i = 0; i < 5; i++)
{
    y[i] = 1;
}
```

y		0x88681140
x[0]		0x88681144
x[1]		0x88681148
x[2]		0x8868114c
x[3]		0x88681150
x[4]		0x88681154
		0x88681158

y	0x88681144	0x88681140
x[0]	1	0x88681144
x[1]	1	0x88681148
x[2]	1	0x8868114c
x[3]	1	0x88681150
x[4]	1	0x88681154
	?	0x88681158

An array name in expression context decays into the array's starting address (address of zero'th element).

However this would **not** be ok: `x = y;` \Rightarrow No space reserved for x to hold an address.

Any pointer can be used with the array access operator `[]`. \Rightarrow y in the example above can be used like x.

Pointer Arithmetics

- The array access operator [] is really only a shorthand for **pointer arithmetic + dereference**
- These are equivalent in C: $x[i] == *(x + i)$

Pointer Arithmetics

- The array access operator [] is really only a shorthand for **pointer arithmetic + dereference**
 - These are equivalent in C: $x[i] == *(x + i)$
-
- The compiler resolves the name of an array to the starting address of the array and adds to it.
 - So the program will happily try to access contents at address $*(x+999999)$, even if array size is much smaller than 999999.
 - Behaviour of exceeding array bounds is "undefined"
 - program might appear to work
 - program might crash (segmentation fault)
 - program might do something apparently random

- Why do pointers have a type, e.g.
`int*` or `char*`?

Pointers Types

- Why do pointers have a type, e.g. `int*` or `char*`?
 - Pointer arithmetic needs to know the size of object that pointer points to so it knows by how much to increment to get to `a[i]`.

y	NULL	0x88681140
x[0]	1	0x88681144
x[1]	1	0x88681148
x[2]	1	0x8868114c
x[3]	1	0x88681150
x[4]	1	0x88681154
	?	0x88681158

Differences between pointers and arrays

```
////////////////////////////////////  
  
int x[5];  
int *y = NULL;  
y = x;  
  
// what is sizeof(x)  
// what is sizeof(y)  
  
func(x);  
func(y);  
  
////////////////////////////////////
```

```
////////////////////////////////////  
  
void func(int x[]) {  
    // what is sizeof(x)  
}  
  
////////////////////////////////////
```

Differences between pointers and arrays

```
////////////////////////////////////
```

```
int x[5];  
int *y = NULL;  
y = x;
```

```
// what is sizeof(x)  
// what is sizeof(y)
```

```
func(x);  
func(y);
```

```
////////////////////////////////////
```

```
////////////////////////////////////
```

```
void func(int x[]) {  
    // what is sizeof(x)  
}
```

```
////////////////////////////////////
```

- Inside main:

$\underbrace{\text{sizeof}(y)}_{\text{size of an address (4 or 8 bytes)}} \neq \underbrace{\text{sizeof}(x)}_{5 * \text{size of an int (typically 20 bytes)}}$

Differences between pointers and arrays

```
////////////////////////////////////
```

```
int x[5];  
int *y = NULL;  
y = x;
```

```
// what is sizeof(x)  
// what is sizeof(y)
```

```
func(x);  
func(y);
```

```
////////////////////////////////////
```

```
////////////////////////////////////
```

```
void func(int x[]) {  
    // what is sizeof(x)  
}
```

```
////////////////////////////////////
```

- Inside main:

$\underbrace{\text{sizeof}(y)}_{\text{size of an address (4 or 8 bytes)}} \neq \underbrace{\text{sizeof}(x)}_{5 * \text{size of an int (typically 20 bytes)}}$

- Inside void func(int x[]):
 sizeof(x) == size of an address

Memory Model

Memory Model

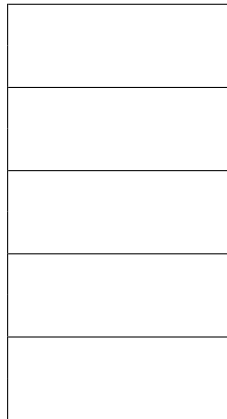
- The memory for a process (a running program) is called its **address space**.
- Memory is just a sequence of bytes.
- A memory location (a byte) is identified by an *address*.

Memory Model

- The memory for a process (a running program) is called its **address space**.
- Memory is just a sequence of bytes.
- A memory location (a byte) is identified by an *address*.
- In A48, you learned about the “memory model”... that was an “incomplete” picture...

Logical address

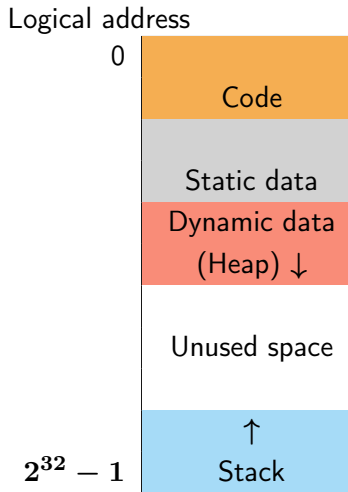
0



$2^{32} - 1$

The Address Space

- **Static data** Space for the *evil* **global** variables and variables declared as **static**
- **Dynamic data (Heap)** Space for dynamically allocated data structures (`malloc`, `calloc`).
- **Stack** Space for variables created in function calls: a function's parameters and a function's local variables



Examples

Logical address

0x

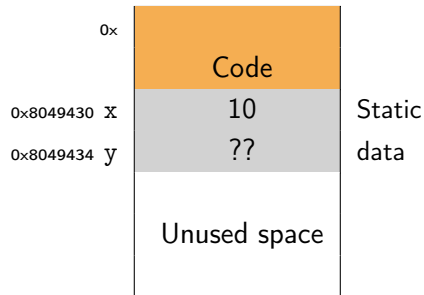
Code

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     return p * q + j;
7 }
8
9 int main() {
10
11     int i = x;
12     y = f(i, i);
13
14     return 0;
15 }
```

Examples

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     return p * q + j;
7 }
8
9 int main() {
10
11     int i = x;
12     y = f(i, i);
13
14     return 0;
15 }
```

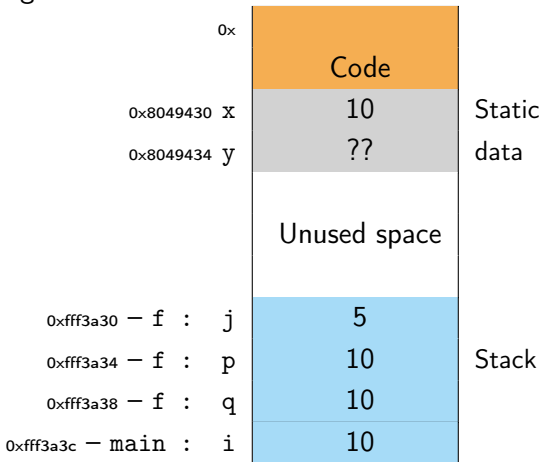
Logical address



Examples

Logical address

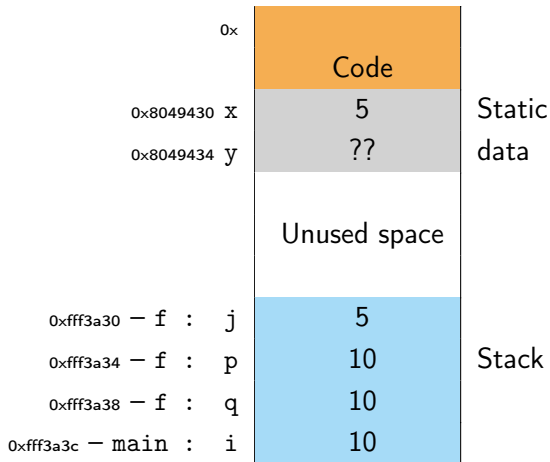
```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     return p * q + j;
7 }
8
9 int main() {
10
11     int i = x;
12     y = f(i, i);
13
14     return 0;
15 }
```



Examples – Global Variables

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     x = 5;
7     return p * q + j;
8 }
9
10 int main() {
11     int i = x;
12     y = f(i, i);
13
14     printf("x=%d, y=%d\n", x, y);
15     return 0;
16 }
17 }
```

If `f()` were to modify `x` or `y`, would this change be permanent?

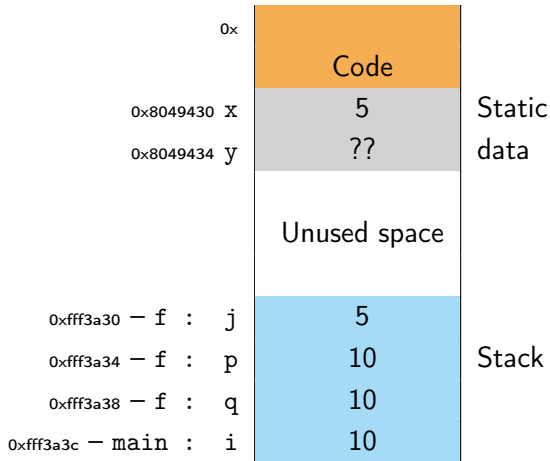


Examples – Global Variables

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     x = 5;
7     return p * q + j;
8 }
9
10 int main() {
11     int i = x;
12     y = f(i, i);
13
14     printf("x=%d, y=%d\n", x, y);
15     return 0;
16 }
17 }
```

If `f()` were to modify `x` or `y`, would this change be permanent?

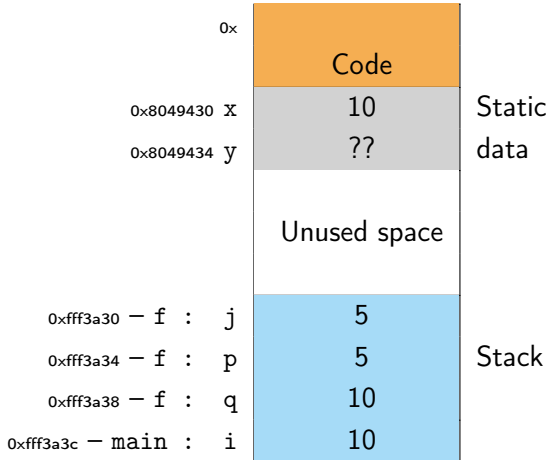
Evil global variables!



Examples – Passing Function Arguments

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     p = 5;
7     return p * q + j;
8 }
9
10 int main() {
11
12     int i = x;
13     y = f(i, i);
14
15     return 0;
16 }
```

If `f()` were to modify `p` or `q`, will that change the value of `main`'s `int i`?

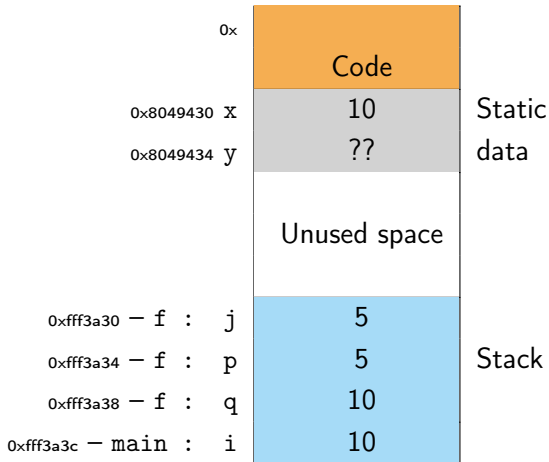


Examples – Passing Function Arguments

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     p = 5;
7     return p * q + j;
8 }
9
10 int main() {
11
12     int i = x;
13     y = f(i, i);
14
15     return 0;
16 }
```

If `f()` were to modify `p` or `q`, will that change the value of `main`'s `int i`?

C passes basic data structures **by value**

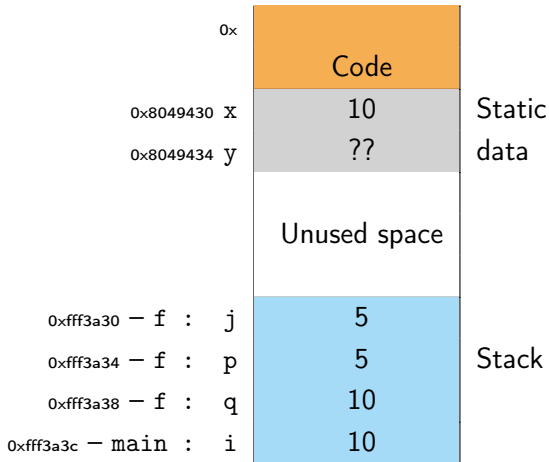


Examples – Passing Function Arguments

```
1 int x = 10;
2 int y;
3
4 int f(int p, int q) {
5     int j = 5;
6     p = 5;
7     return p * q + j;
8 }
9
10 int main() {
11
12     int i = x;
13     y = f(i, i);
14
15     return 0;
16 }
```

If `f()` were to modify `p` or `q`, will that change the value of `main`'s `int i`?

C passes basic data structures **by value** ⇒ What change is needed for `f()` to change one of `main`'s variables?



Examples – Passing Function Arguments


```
1 int x = 10;
2 int y;
3
4 void f(int *p, int q) {
5     *p = 5;
6 }
7
8 int main() {
9
10     int i = x;
11     f(&i, i);
12     return 0;
13 }
```

Changes

f(&i,i);

void f(int *p, q)

*p = 5;



Examples – Passing Function Arguments

```
1 int x = 10;
2 int y;
3
4 void f(int *p, int q) {
5     *p = 5;
6 }
7
8 int main() {
9
10     int i = x;
11     f(&i, i);
12     return 0;
13 }
```

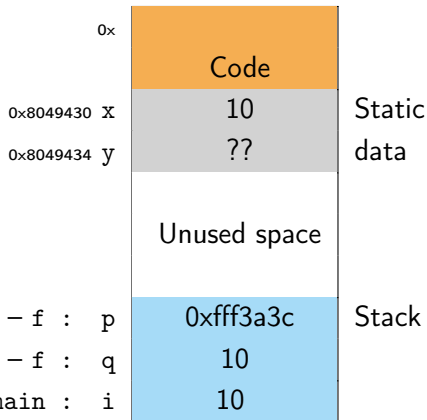
When calling f()

Changes

f(&i,i);

void f(int *p, q)

*p = 5;



Examples – Passing Function Arguments

```
1 int x = 10;
2 int y;
3
4 void f(int *p, int q) {
5     *p = 5;
6 }
7
8 int main() {
9
10     int i = x;
11     f(&i, i);
12     return 0;
13 }
```

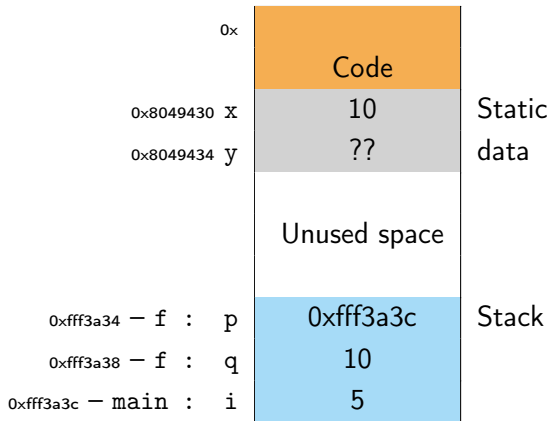
After f() returns

Changes

f(&i,i);

void f(int *p, q)

*p = 5;



The Address Space, continued...

- **Static data**

.data = global or static variables with predefined values that can be modified

BSS segment = global/static without predefined values

- **Dynamic data (Heap)**

- **Stack**

- **Code/Text**

Logical address

0

Code

Static data

Dynamic data
(Heap) ↓

Unused space

$2^{32} - 1$

↑

Stack

Dynamic Memory Management

Let's write our own concat function:

```
1 char *concat(const char *s1,  
2             const char *s2) {  
3     char result[70];  
4  
5     strcpy(result, s1);  
6     strcat(result, s2);  
7  
8     return result;  
9 }
```

Any problems with this implementation?

Dynamic Memory Management

Let's write our own concat function:

```
1 char *concat(const char *s1,  
2             const char *s2) {  
3     char result[70];  
4  
5     strcpy(result, s1);  
6     strcat(result, s2);  
7  
8     return result;  
9 }
```

Any problems with this implementation?

- strncpy & strncat would be safer..
- But what else?

Dynamic Memory Management

Let's write our own concat function:

```
1 char *concat(const char *s1,  
2             const char *s2) {  
3     char result[70];  
4  
5     strcpy(result, s1);  
6     strcat(result, s2);  
7  
8     return result;  
9 }
```

```
1 char *concat(const char *s1, const char *s2)  
2 {  
3     // temp local variable to store result  
4     char *result;  
5  
6     // allocating enough memory  
7     result = malloc(strlen(s1) + strlen(s2)  
8                 + 1);  
9     if (result == NULL) {  
10         printf ( Error: malloc failed\n  
11                 );  
12         exit(1);  
13     }  
14     //  
15     strcpy (result, s1);  
16     strcat (result, s2);  
17     return result;  
18 }
```

The Address Space, continued...

- **Static data**

- **Dynamic data (Heap)**

Space for dynamically allocated data structures (`malloc`, `calloc`).

- Memory allocated here will never be released/freed automatically by the system.
- Completely under programmers control.
- Memory here can be allocated at any time during the run of a program

- **Stack**

- **Code/Text**

Logical address

0

Code

Static data

Dynamic data
(Heap) ↓

Unused space

↑

Stack

$2^{32} - 1$

Differences between `char` pointers and `char` arrays

```
// char array
char a[] = "array";

// char pointer
// AKA "string literal"
char *p = "pointer";
```

```
p[0] = 'z';    // Illegal!
a[0] = 'z';    // OK!
```

"pointer" is stored in **read-only memory**.

No other space is reserved for `p`, except to store a memory address.

The *Complete* Memory Model

- **Static data**

Space for the *evil* **global** variables and variables declared as **static**

- **Dynamic data (Heap)**

Space for dynamically allocated data structures (malloc, calloc).

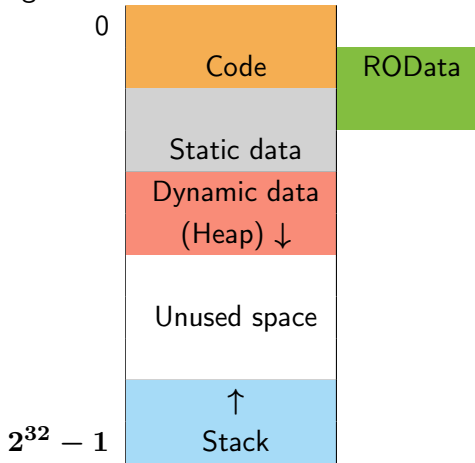
- **Stack**

Space for variables created in function calls: a function's parameters and a function's local variables

- **Code/Text**

string literals: ROData (read-only data)
code/text, or static data depending on platform

Logical address



Passing Arrays as Parameters

Suppose we want to write a function to sum an array of ints.

```
int main()
{
    int i[3] = {10, 9, 8};

    printf("sum is %d\n", sum(i)); /*??*/

    return 0;
}

int sum( /* What goes here? */ ) {
}
```

What is being passed to the function is the name of the array which *decays* to a pointer to the first element – a pointer of type `int`.

Passing Arrays as Parameters

Suppose we want to write a function to sum an array of ints.

```
int main()
{
    int i[3] = {10, 9, 8};

    printf("sum is %d\n", sum(i)); /*??*/

    return 0;
}

int sum( /* What goes here? */ ) {
}
```

```
int sum( int *a ) {
    int i, s = 0;

    for(i = 0; i < ??; i++)
        s += a[i]; /* this is
                    legal */

    return s;
}
```

Passing Arrays as Parameters

Suppose we want to write a function to sum an array of ints.

```
int main()
{
    int i[3] = {10, 9, 8};

    printf("sum is %d\n", sum(i)); /*??*/

    return 0;
}

int sum( /* What goes here? */ ) {
}
```

```
int sum( int *a ) {
    int i, s = 0;

    for(i = 0; i < ??; i++)
        s += a[i]; /* this is
                    legal */

    return s;
}
```

- How do we know how big the array is?
- Pass in the *size* of the array as another parameter.

Array Parameters

- `int sum(int *a, int size)`
- Also valid C code is:
`int sum(int a[], int size)`

Array Parameters

- `int sum(int *a, int size)`
- Also valid C code is:
`int sum(int a[], int size)`
- However, many reasons against using this form:
 - We are really passing a pointer-to-int **not** an array.
 - We still don't know how big the array is.
 - Outside of a *formal parameter declaration* `int a[]`; is **illegal**
- `int a`; and `int a[10]`; are completely different things

- Arrays are **not** pointers
- “Equivalence” of pointers and arrays

Decay: An *lvalue* of type *array-of-T* which appears in an expression **decays** (with three exceptions) into a pointer to its first element; the type of the resultant pointer is pointer-to-T.

(The exceptions are when the array is the operand of a `sizeof` or `&` operator, or is a *literal string* initializer for a character array.)

- Array and pointer declarations are interchangeable as function formal parameters
- Arrays automatically allocate space, but can't be relocated or resized.
Pointers must be explicitly assigned to point to allocated space (perhaps using `malloc`), but can be reassigned (i.e. pointed at different objects) at will, and have many other uses besides serving as the base of blocks of memory.

- Due to the so-called equivalence of arrays and pointers, arrays and pointers often seem interchangeable, and in particular a pointer to a block of memory assigned by `malloc` is frequently treated (and can be referenced using `[]` exactly) as if it were a true array.